

ex_nn

December 3, 2021

1 Lena Feiler - i6246119

2 Machine Learning: Artificial Neural Networks

Instructions _____

This file contains code that helps you get started. You will need to complete the following functions

- predict.m
- sigmoidGradient.m
- randInitializeWeights.m
- nnCostFunction.m

For this exercise, you will not need to change any code in this file, or any other files other than those mentioned above.

2.1 Import the required packages

```
[1]: import scipy.io
import numpy as np

from predict import predict
from displayData import displayData
from sigmoidGradient import sigmoidGradient
from randInitializeWeights import randInitializeWeights
from nnCostFunction import nnCostFunction
from checkNNGradients import checkNNGradients
from fmincg import fmincg
```

2.2 Setup the parameters you will use for this exercise

```
[2]: input_layer_size = 400;      # 20x20 Input Images of Digits
hidden_layer_size = 25;          # 25 hidden units
num_labels = 10;                 # 10 labels, from 0 to 9
                                   # (note that we have mapped "0" to label 9 to follow
                                   # the same structure used in the MatLab version)
```

3 ===== Part 1: Loading and Visualizing Data =====

We start the exercise by first loading and visualizing the dataset. You will be working with a dataset that contains handwritten digits.

3.1 Load Training Data

```
[3]: print('Loading and Visualizing Data ...')

mat = scipy.io.loadmat('digitdata.mat')
X = mat['X']
y = mat['y']
y = np.squeeze(y)
m, _ = np.shape(X)


# Randomly select 100 data points to display
sel = np.random.choice(range(X.shape[0]), 100)
sel = X[sel,:]

displayData(sel)
```

Loading and Visualizing Data ...

/Users/lena/DKE/Year_2/period_2.2/machine_learning/Lab5/displayData.py:8:
MatplotlibDeprecationWarning: Passing non-integers as three-element position
specification is deprecated since 3.3 and will be removed two minor releases
later.

```
ax = fig.add_subplot(np.sqrt(data.shape[0]), np.sqrt(data.shape[0]), i+1)
```



4 ===== Part 2: Loading Parameters =====

In this part of the exercise, we load some pre-initialized neural network parameters.

```
[4]: print('Loading Saved Neural Network Parameters ...')

# Load the weights into variables Theta1 and Theta2
mat = scipy.io.loadmat('debugweights.mat');

# Unroll parameters
Theta1 = mat['Theta1']
Theta1_1d = np.reshape(Theta1, Theta1.size, order='F')
Theta2 = mat['Theta2']
Theta2_1d = np.reshape(Theta2, Theta2.size, order='F')

nn_params = np.hstack((Theta1_1d, Theta2_1d))
```

Loading Saved Neural Network Parameters ...

5 ===== Part 3: Implement Predict =====

After training the neural network, we would like to use it to predict the labels. You will now implement the “predict” function to use the neural network to predict the labels of the training set. This lets you compute the training set accuracy.

```
[5]: pred = predict(Theta1, Theta2, X);
print('Training Set Accuracy: ', (pred == y).mean()*100)
```

Training Set Accuracy: 97.52

5.1 Testing (you can skip this block)

To give you an idea of the network’s output, you can also run through the examples one at the a time to see what it is predicting. Run the code in the following block to view examples.

NOTE: to avoid the printing of all the sample instances, you can replace $range(m)$ with a small number

```
[ ]: # Randomly permute examples
rp = np.random.permutation(m)

for i in range(m):
    print(i)
    # Display
```

```

print('Displaying Example Image')
tmp = np.transpose(np.expand_dims(X[rp[i], :], axis=1))
displayData(tmp)

pred = predict(Theta1, Theta2, tmp)
print('Neural Network Prediction: ', pred, '(digit ', pred%10, ')')

```

6 ===== Part 4: Sigmoid Gradient =====

Before you start implementing backpropagation, you will first implement the gradient for the sigmoid function. You should complete the code in the sigmoidGradient.m file.

```

[6]: print('Evaluating sigmoid gradient...')
example = np.array([-15, -1, -0.5, 0, 0.5, 1, 15])
g = sigmoidGradient(example)
print('Sigmoid gradient evaluated at', example, ':')
print(g)

```

```

Evaluating sigmoid gradient...
Sigmoid gradient evaluated at [-15.  -1.  -0.5  0.   0.5  1.  15.] :
[3.05902133e-07  1.96611933e-01  2.35003712e-01  2.50000000e-01
 2.35003712e-01  1.96611933e-01  3.05902133e-07]

```

7 ===== Part 5: Initializing Parameters =====

To learn a two layer neural network that classifies digits. You will start by implementing a function to initialize the weights of the neural network (randInitializeWeights.py)

```

[7]: print('Initializing Neural Network Parameters ...')

initial_Theta1 = randInitializeWeights(input_layer_size, hidden_layer_size)
initial_Theta2 = randInitializeWeights(hidden_layer_size, num_labels)

# Unroll parameters
initial_Theta1 = np.reshape(initial_Theta1, initial_Theta1.size, order='F')
initial_Theta2 = np.reshape(initial_Theta2, initial_Theta2.size, order='F')
initial_nn_params = np.hstack((initial_Theta1, initial_Theta2))
print(initial_nn_params)

```

```

Initializing Neural Network Parameters ...
[-0.09458821 -0.08262094 -0.04324267 ...  0.07132938 -0.03409781
 -0.10495504]

```

8 ===== Part 6: Implement Backpropagation =====

Now you will implement the backpropagation algorithm for the neural network. You should add code to `nnCostFunction.m` to return the partial derivatives of the parameters.

```
[8]: print('Checking Backpropagation...')  
  
# Check gradients by running checkNNGradients  
checkNNGradients()
```

Checking Backpropagation...

```
[[ -9.27825235e-03  
   [ 8.89911959e-03  
   [-8.36010761e-03  
   [ 7.62813551e-03  
   [-6.74798369e-03  
   [-3.04978931e-06  
   [ 1.42869450e-05  
   [-2.59383093e-05  
   [ 3.69883213e-05  
   [-4.68759787e-05  
   [-1.75060084e-04  
   [ 2.33146356e-04  
   [-2.87468729e-04  
   [ 3.35320347e-04  
   [-3.76215588e-04  
   [-9.62660640e-05  
   [ 1.17982668e-04  
   [-1.37149705e-04  
   [ 1.53247079e-04  
   [-1.66560297e-04  
   [ 3.14544970e-01  
   [ 1.11056588e-01  
   [ 9.74006970e-02  
   [ 1.64090819e-01  
   [ 5.75736494e-02  
   [ 5.04575855e-02  
   [ 1.64567932e-01  
   [ 5.77867378e-02  
   [ 5.07530173e-02  
   [ 1.58339334e-01  
   [ 5.59235296e-02  
   [ 4.91620841e-02  
   [ 1.51127527e-01  
   [ 5.36967009e-02  
   [ 4.71456249e-02  
   [ 1.49568335e-01]
```

```

[ 5.31542052e-02]
[ 4.65597186e-02]] [[-9.27825236e-03]
[ 8.89911960e-03]
[-8.36010762e-03]
[ 7.62813551e-03]
[-6.74798370e-03]
[-3.04978914e-06]
[ 1.42869443e-05]
[-2.59383100e-05]
[ 3.69883234e-05]
[-4.68759769e-05]
[-1.75060082e-04]
[ 2.33146357e-04]
[-2.87468729e-04]
[ 3.35320347e-04]
[-3.76215587e-04]
[-9.62660620e-05]
[ 1.17982666e-04]
[-1.37149706e-04]
[ 1.53247082e-04]
[-1.66560294e-04]
[ 3.14544970e-01]
[ 1.11056588e-01]
[ 9.74006970e-02]
[ 1.64090819e-01]
[ 5.75736493e-02]
[ 5.04575855e-02]
[ 1.64567932e-01]
[ 5.77867378e-02]
[ 5.07530173e-02]
[ 1.58339334e-01]
[ 5.59235296e-02]
[ 4.91620841e-02]
[ 1.51127527e-01]
[ 5.36967009e-02]
[ 4.71456249e-02]
[ 1.49568335e-01]
[ 5.31542052e-02]
[ 4.65597186e-02]]

```

The above two columns you get should be very similar.
 (Left-Numerical Gradient, Right-(Your) Analytical Gradient)

If your backpropagation implementation is correct, then
 the relative difference will be small (less than 1e-9).

Relative Difference: 2.2957544655995076e-11

9 ===== Part 7: Implement Regularization =====

Once your backpropagation implementation is correct, you should now continue to implement the regularization gradient.

```
[9]: print('Checking Backpropagation (w/ Regularization) ... ')

    ## Check gradients by running checkNNGradients
    lambda_value = 3
    checkNNGradients(lambda_value)

    # Also output the costFunction debugging values
    debug_J = nnCostFunction(nn_params, input_layer_size, hidden_layer_size,
                             num_labels, X, y, lambda_value)

    print('Cost at (fixed) debugging parameters (w/ lambda = 10): ', debug_J[0][0],
          '(this value should be about 0.576051)')
```

Checking Backpropagation (w/ Regularization) ...

```
[[ -9.27825235e-03]
 [  8.89911959e-03]
 [-8.36010761e-03]
 [  7.62813551e-03]
 [-6.74798369e-03]
 [-1.67679797e-02]
 [  3.94334829e-02]
 [  5.93355565e-02]
 [  2.47640974e-02]
 [-3.26881426e-02]
 [-6.01744725e-02]
 [-3.19612287e-02]
 [  2.49225535e-02]
 [  5.97717617e-02]
 [  3.86410548e-02]
 [-1.73704651e-02]
 [-5.75658668e-02]
 [-4.51963845e-02]
 [  9.14587966e-03]
 [  5.46101547e-02]
 [  3.14544970e-01]
 [  1.11056588e-01]
 [  9.74006970e-02]
 [  1.18682669e-01]
 [  3.81928689e-05]
 [  3.36926556e-02]
 [  2.03987128e-01]
 [  1.17148233e-01]
```

```

[ 7.54801264e-02]
[ 1.25698067e-01]
[-4.07588279e-03]
[ 1.69677090e-02]
[ 1.76337550e-01]
[ 1.13133142e-01]
[ 8.61628953e-02]
[ 1.32294136e-01]
[-4.52964427e-03]
[ 1.50048382e-03]] [[-9.27825236e-03]
[ 8.89911960e-03]
[-8.36010762e-03]
[ 7.62813551e-03]
[-6.74798370e-03]
[-1.67679797e-02]
[ 3.94334829e-02]
[ 5.93355565e-02]
[ 2.47640974e-02]
[-3.26881426e-02]
[-6.01744725e-02]
[-3.19612287e-02]
[ 2.49225535e-02]
[ 5.97717617e-02]
[ 3.86410548e-02]
[-1.73704651e-02]
[-5.75658668e-02]
[-4.51963845e-02]
[ 9.14587966e-03]
[ 5.46101547e-02]
[ 3.14544970e-01]
[ 1.11056588e-01]
[ 9.74006970e-02]
[ 1.18682669e-01]
[ 3.81928696e-05]
[ 3.36926556e-02]
[ 2.03987128e-01]
[ 1.17148233e-01]
[ 7.54801264e-02]
[ 1.25698067e-01]
[-4.07588279e-03]
[ 1.69677090e-02]
[ 1.76337550e-01]
[ 1.13133142e-01]
[ 8.61628953e-02]
[ 1.32294136e-01]
[-4.52964427e-03]
[ 1.50048382e-03]]

```

The above two columns you get should be very similar.

(Left-Numerical Gradient, Right-(Your) Analytical Gradient)

If your backpropagation implementation is correct, then
the relative difference will be small (less than $1e-9$).

Relative Difference: 2.2006043191433586e-11

Cost at (fixed) debugging parameters (w/ $\lambda = 10$): 0.5760512469501331 (this
value should be about 0.576051)

10 ===== Part 8: Training NN =====

You have now implemented all the code necessary to train a neural network. To train your neural network, we will now use “fmincg”, which is a function which works similarly to “fminunc”. Recall that these advanced optimizers are able to train our cost functions efficiently as long as we provide them with the gradient computations.

```
[10]: print('Training Neural Network...')

# After you have completed the assignment, change the MaxIter to a larger
# value to see how more training helps.
MaxIter = 150

# You should also try different values of lambda
lambda_value = 1

# Create "short hand" for the cost function to be minimized
y = np.expand_dims(y, axis=1)

costFunction = lambda p : nnCostFunction(p, input_layer_size, hidden_layer_size,
                                         num_labels, X, y, lambda_value)

# Now, costFunction is a function that takes in only one argument (the
# neural network parameters)
[nn_params, cost] = fmincg(costFunction, initial_nn_params, MaxIter)

# Obtain Theta1 and Theta2 back from nn_params
Theta1 = np.reshape(nn_params[0:hidden_layer_size * (input_layer_size + 1)],
                    (hidden_layer_size, (input_layer_size + 1)),
                    order='F')
Theta2 = np.reshape(nn_params[((hidden_layer_size * (input_layer_size + 1))):],
                    (num_labels, (hidden_layer_size + 1)), order='F')
```

Training Neural Network...

Iteration 1 | Cost: [3.28439445]

Iteration 2 | Cost: [3.24307806]

Iteration	3	Cost:	[3.20108698]
Iteration	4	Cost:	[2.88219067]
Iteration	5	Cost:	[2.5838405]
Iteration	6	Cost:	[2.40483167]
Iteration	7	Cost:	[2.11419278]
Iteration	8	Cost:	[1.85711249]
Iteration	9	Cost:	[1.71349158]
Iteration	10	Cost:	[1.53586494]
Iteration	11	Cost:	[1.46577259]
Iteration	12	Cost:	[1.43507973]
Iteration	13	Cost:	[1.31211977]
Iteration	14	Cost:	[1.27110142]
Iteration	15	Cost:	[1.23787513]
Iteration	16	Cost:	[1.1909416]
Iteration	17	Cost:	[1.13230706]
Iteration	18	Cost:	[1.10533307]
Iteration	19	Cost:	[1.07781607]
Iteration	20	Cost:	[1.01960256]
Iteration	21	Cost:	[0.94854419]
Iteration	22	Cost:	[0.91385642]
Iteration	23	Cost:	[0.89375593]
Iteration	24	Cost:	[0.84258204]
Iteration	25	Cost:	[0.82934597]
Iteration	26	Cost:	[0.81361549]
Iteration	27	Cost:	[0.77770115]
Iteration	28	Cost:	[0.73481581]
Iteration	29	Cost:	[0.69344115]
Iteration	30	Cost:	[0.66212686]
Iteration	31	Cost:	[0.64897138]
Iteration	32	Cost:	[0.63451819]
Iteration	33	Cost:	[0.62062928]
Iteration	34	Cost:	[0.61105205]
Iteration	35	Cost:	[0.60394273]
Iteration	36	Cost:	[0.58742808]
Iteration	37	Cost:	[0.57452468]
Iteration	38	Cost:	[0.56952945]
Iteration	39	Cost:	[0.56368648]
Iteration	40	Cost:	[0.55752858]
Iteration	41	Cost:	[0.55443359]
Iteration	42	Cost:	[0.550767]
Iteration	43	Cost:	[0.53937685]
Iteration	44	Cost:	[0.52041994]
Iteration	45	Cost:	[0.50907771]
Iteration	46	Cost:	[0.49821091]
Iteration	47	Cost:	[0.48588888]
Iteration	48	Cost:	[0.47916008]
Iteration	49	Cost:	[0.47746599]
Iteration	50	Cost:	[0.47027926]

Iteration	51	Cost:	[0.4673045]
Iteration	52	Cost:	[0.46576225]
Iteration	53	Cost:	[0.46520005]
Iteration	54	Cost:	[0.46278354]
Iteration	55	Cost:	[0.46105291]
Iteration	56	Cost:	[0.46003879]
Iteration	57	Cost:	[0.45903774]
Iteration	58	Cost:	[0.45577482]
Iteration	59	Cost:	[0.45194366]
Iteration	60	Cost:	[0.45112513]
Iteration	61	Cost:	[0.44878775]
Iteration	62	Cost:	[0.43973194]
Iteration	63	Cost:	[0.43607229]
Iteration	64	Cost:	[0.43369701]
Iteration	65	Cost:	[0.43247161]
Iteration	66	Cost:	[0.43064263]
Iteration	67	Cost:	[0.4290564]
Iteration	68	Cost:	[0.42836698]
Iteration	69	Cost:	[0.42747387]
Iteration	70	Cost:	[0.4268125]
Iteration	71	Cost:	[0.42459993]
Iteration	72	Cost:	[0.42323659]
Iteration	73	Cost:	[0.42008503]
Iteration	74	Cost:	[0.41801322]
Iteration	75	Cost:	[0.41560441]
Iteration	76	Cost:	[0.41322673]
Iteration	77	Cost:	[0.41148567]
Iteration	78	Cost:	[0.41010694]
Iteration	79	Cost:	[0.40746177]
Iteration	80	Cost:	[0.40438447]
Iteration	81	Cost:	[0.40350454]
Iteration	82	Cost:	[0.40298928]
Iteration	83	Cost:	[0.40288494]
Iteration	84	Cost:	[0.40249783]
Iteration	85	Cost:	[0.40204832]
Iteration	86	Cost:	[0.40168254]
Iteration	87	Cost:	[0.40041093]
Iteration	88	Cost:	[0.39834516]
Iteration	89	Cost:	[0.39302062]
Iteration	90	Cost:	[0.38847072]
Iteration	91	Cost:	[0.38694872]
Iteration	92	Cost:	[0.38494207]
Iteration	93	Cost:	[0.38315387]
Iteration	94	Cost:	[0.38099824]
Iteration	95	Cost:	[0.3790177]
Iteration	96	Cost:	[0.3775966]
Iteration	97	Cost:	[0.37645034]
Iteration	98	Cost:	[0.37495215]

Iteration	99	Cost:	[0.3742671]
Iteration	100	Cost:	[0.37389859]
Iteration	101	Cost:	[0.37360405]
Iteration	102	Cost:	[0.37312382]
Iteration	103	Cost:	[0.37237434]
Iteration	104	Cost:	[0.37194321]
Iteration	105	Cost:	[0.37173018]
Iteration	106	Cost:	[0.37118032]
Iteration	107	Cost:	[0.37082975]
Iteration	108	Cost:	[0.3707602]
Iteration	109	Cost:	[0.37049161]
Iteration	110	Cost:	[0.37038189]
Iteration	111	Cost:	[0.37035186]
Iteration	112	Cost:	[0.37014281]
Iteration	113	Cost:	[0.36988736]
Iteration	114	Cost:	[0.36965214]
Iteration	115	Cost:	[0.36889956]
Iteration	116	Cost:	[0.36699026]
Iteration	117	Cost:	[0.36205054]
Iteration	118	Cost:	[0.35635538]
Iteration	119	Cost:	[0.3548609]
Iteration	120	Cost:	[0.3544416]
Iteration	121	Cost:	[0.3533465]
Iteration	122	Cost:	[0.35234534]
Iteration	123	Cost:	[0.35182596]
Iteration	124	Cost:	[0.35149403]
Iteration	125	Cost:	[0.35120181]
Iteration	126	Cost:	[0.35087614]
Iteration	127	Cost:	[0.35041806]
Iteration	128	Cost:	[0.35023378]
Iteration	129	Cost:	[0.35017133]
Iteration	130	Cost:	[0.34998164]
Iteration	131	Cost:	[0.34989407]
Iteration	132	Cost:	[0.34979186]
Iteration	133	Cost:	[0.34954061]
Iteration	134	Cost:	[0.34920649]
Iteration	135	Cost:	[0.34883394]
Iteration	136	Cost:	[0.34871074]
Iteration	137	Cost:	[0.34861334]
Iteration	138	Cost:	[0.34849444]
Iteration	139	Cost:	[0.34823424]
Iteration	140	Cost:	[0.3476669]
Iteration	141	Cost:	[0.34708856]
Iteration	142	Cost:	[0.34641996]
Iteration	143	Cost:	[0.34518428]
Iteration	144	Cost:	[0.34467471]
Iteration	145	Cost:	[0.34395803]
Iteration	146	Cost:	[0.34329639]

```
Iteration 147 | Cost: [0.34304632]
Iteration 148 | Cost: [0.3429067]
Iteration 149 | Cost: [0.34275092]
Iteration 150 | Cost: [0.34256184]
```

11 ===== Part 9: Visualize Weights =====

You can now “visualize” what the neural network is learning by displaying the hidden units to see what features they are capturing in the data.

```
[11]: print('\nVisualizing Neural Network... \n')

displayData(Theta1[:, 1:])
```

Visualizing Neural Network...



12 ===== Part 10: Predicting with learned weights =====

After training the neural network, we would like to use it to predict the labels. The already implemented “predict” function is used by neural network to predict the labels of the training set. This lets you compute the training set accuracy.

```
[12]: pred = predict(Theta1, Theta2, X)
pred = np.expand_dims(pred,axis=1)
print('Training Set Accuracy: ', (pred == y).mean()*100)
```

Training Set Accuracy: 98.9

```
[ ]:
```

13 Evaluation/ Report

13.1 Description of the network (e.g. number of layers, number of neurons per layer, etc.)

The network, as described in the assignment, has 3 layers, the input and output layer and one hidden layer. The input layer has 401 units, where 1 is the bias node. The hidden layer has 26 neurons (again, 1 bias node) and the output layer has 10. The output units correspond to the 10 digit classes.

13.2 Impact of specific parameters such as , number of iterations, weight initialization, etc.

- Impact of lambda:

Lambda is our regularization rate. It is used in the nnCostFunction in the regularization process. It helps us limit how much regularisation we allow. Increasing its value effect of the regularization. If its value is increased too much, the model will become more general and can then happen to underfit the given data. On the other hand, if it is too low, we risk overfitting the data, since our model becomes more complex, since we allow the model to learn more about the training data.

- Impact of weight initialization:

If both bias and weights are initialized to 0, all neurons will learn the same, since the output of the neurons and the backpropagation of the error, respectively, are the same. Otherwise, if all values are initialized randomly, the vanishing gradient can occur, since the values can be very high and therefore the value of the activation function is close to 0. Therefore, we know that the initialization of the weights can have a very high impact.

- Impact of max iterations:

Increasing the number of allowed iterations (maxIterations), can help decrease the error and increase the accuracy of the model. Therefore, by increasing the number of iterations, the error will get closer to 0, but does not necessarily have to converge.

- Impact of weight epsilon:

Epsilon is the range of the weight initialization. If epsilon is large, the random weights may take values that are very different from one another, which can also create outliers. On the other hand, if epsilon is small, the weights will be more similar to one another, which can make learning more efficient.

13.3 How does the regularization affect the training of your ANN?

Regularization is used to reduce the degree of freedom used to update and hence only make small modifications to the applied learning algorithm, in order to improve the model's generalization. Therefore, it may improve the performance on unseen data, since it limits overfitting of the model.

Running the code in part 7 with regularization and without regularization, gives us the following values:

with Regularization: Relative Difference: 2.2006043191433586e-11

without Regularization: Relative Difference: 0.22177454344825828

This shows that regularization decreased the relative difference between the given numerical Gradient, and the acquired Analytical Gradient, produced by the ANN. Therefore, the produced output has become more accurate.

13.4 Did you manage to improve the initial results (using values in debugweights.mat)? Which was your best result? How did you configure the system? How could you improve them even more?

At the beginning of 'Part 3: Implement Predict', we check the accuracy of the model using the values in debugweights.mat, using only the feedforward implementation. This gives us the following accuracy:

Training Set Accuracy: 97.52

After implementing the backpropagation, the accuracy is again checked in 'Part 10: Predicting with learned weights'. This gives us:

Training Set Accuracy: 98.66

This means, that our trained weights perform better than the initial given weights. Therefore, implementing the backpropagation leads to an increase in accuracy.

To Further improve, we could add more neurons in the hidden layer, or more layers. However, this can also lead to the model overfitting and losing generalization.

13.5 Imagine that you want to use a similar solution to classify 50x50 pixel grayscale images containing letters (consider an alphabet with 26 letters). Which changes would you need in the current code in order to implement this classification task?

To apply the current code to the new problem, we would only need to redefine the input values in the section **Setup the parameters you will use for this exercise**.

In our current network we have a 20x20 grayscale image. So, to change to a 50x50 pixel grayscale image, we would have to change the **input layer** to have **2501 units** (2500 + 1 bias node).

Furthermore, since we now have 26 letters we want to classify, instead of 10 numbers, the **output layer** should be adjusted to have **27 neurons** (26 for the letters and 1 bias unit).

For the **hidden layer**, no explicit number is specified. However, it could be considered to add more layers. This has a chance of improving the accuracy, although this is not guaranteed, and might even cause the accuracy to decrease, if too many layers are added, since the model will end

up overfitting. However, if we would add more hidden layers, we would need to change more parts of the code, therefore, we will continue working with 1 layer. For the number of neurons in the hidden layer(s), the number is also not defined. Since the problems' complexity has increased, we should also increase the number of hidden neurons, and for example use 501 neurons. This value can be adjusted.

So, our changes would be: `> input_layer_size = 2500;`

```
hidden_layer_size = 500;
```

```
num_labels = 26;
```

13.6 Change the value of the variable `show_examples` (in the python version, run the relevant block in the Jupyter one) in `ex_nn`, which information is provided? Did you get the expected information? Is anything unexpected there?

`Show_examples` is a boolean variable. Therefore, we change its value from `False` to `True`. When executing the corresponding part in jupyter notebook, we are shown, for example:

21

Displaying Example Image

Neural Network Prediction: [6] (digit [6])

So, they show the number of the examples that is being displayed, the sentence 'Displaying Example Image', the digit that was detected '[6]', as well as the actual value of the current example '(digit [6])'. Furthermore, the image that is being analyzed is also shown.

13.7 How does your `sigmoidGradient` function work? Which is the return value for different values of `z`? How does it work with the input is a vector and with it is a matrix?

The `sigmoidGradient` method computes the gradient of the sigmoid function evaluated at every given value of `z`. The method returns the sigmoid function of `z`, multiplied with the sigmoid function of `z` subtracted from 1, so: `> g = sigmoid(z) * (1 - sigmoid(z))`

The method is defined to be able to take `z` as a matrix, vector or scalar. Each example can also be seen in the code below. The type of the input value is the same as the type of the output value.

1. If `z` is a scalar, the output value will be also be a scalar.
2. If `z` is a vector, the output value will be also be a vector.
3. If `z` is a matrix, the output value will be also a matrix.

```
[13]: # scalar
sclar = 1
s = sigmoidGradient(1)
print('output of a scalar input: ', s)
print('shape: ', np.shape(s))
```



```

# vector
vector = np.array([-15, -1, -0.5, 0, 0.5, 1, 15])
v = sigmoidGradient(vector)
print('output of a vector input: ', v)
print('shape: ', np.shape(v))

# matrix
matrix = np.array([[4,7,6],[1,2,5],[9,3,8]])
m = sigmoidGradient(matrix)
print('output of a matrix inout: ', m)
print('shape: ', np.shape(m))

```

```

output of a scalar input:  0.19661193324148185
shape:  ()
output of a vector input:  [3.05902133e-07 1.96611933e-01 2.35003712e-01
2.50000000e-01
 2.35003712e-01 1.96611933e-01 3.05902133e-07]
shape:  (7,)
output of a matrix inout:  [[1.76627062e-02 9.10221180e-04 2.46650929e-03]
 [1.96611933e-01 1.04993585e-01 6.64805667e-03]
 [1.23379350e-04 4.51766597e-02 3.35237671e-04]]
shape:  (3, 3)

```

[]: